

Prediction Assignment Writeup

Sabin Shrestha

June 5, 2016

Project Requirement:

The goal of your project is to predict the manner in which they did the exercise. This is the “classe” variable in the training set. You may use any of the other variables to predict with. You should create a report describing how you built your model, how you used cross validation, what you think the expected out of sample error is, and why you made the choices you did. You will also use your prediction model to predict 20 different test cases.

Data: The training data for this project are available here: source: <http://groupware.les.inf.puc-rio.br/har.https://d396qusza40orc.cloudfront.net/predmachlearn/pml-training.csv>

The test data are available here:

<https://d396qusza40orc.cloudfront.net/predmachlearn/pml-testing.csv>

Setting up the environment

```
#install.packages("RGtk2")
#install.packages("rattle")
#install.packages('rpart')

#install.packages('lattice')
#install.packages('ggplot2')
#install.packages('rpart.plot')

library(lattice)
```

```
## Warning: package 'lattice' was built under R version 3.2.5
```

```
library(ggplot2)
```

```
## Warning: package 'ggplot2' was built under R version 3.2.5
```

```
library(caret)
```

```
## Warning: package 'caret' was built under R version 3.2.5
```

```
library(rpart)
```

```
## Warning: package 'rpart' was built under R version 3.2.5
```

```
library(rpart.plot)
```

```
## Warning: package 'rpart.plot' was built under R version 3.2.5
```

```
library(RColorBrewer)
library(rattle)
```

```
## Warning: package 'rattle' was built under R version 3.2.5
```

```
## Rattle: A free graphical interface for data mining with R.
## Version 4.1.0 Copyright (c) 2006-2015 Togaware Pty Ltd.
## Type 'rattle()' to shake, rattle, and roll your data.
```

```
library(randomForest)
```

```
## Warning: package 'randomForest' was built under R version 3.2.5
```

```
## randomForest 4.6-12
```

```
## Type rfNews() to see new features/changes/bug fixes.
```

```
##
## Attaching package: 'randomForest'
```

```
## The following object is masked from 'package:ggplot2':
##
##      margin
```

Data Loading

```
pmlTraining<-read.csv("D:/Data_Science/Practical-Machine-Learning/Data/pml-training.csv", header=T, na.st
pmlTesting<-read.csv("D:/Data_Science/Practical-Machine-Learning/Data/pml-testing.csv", header=T, na.st
dim(pmlTraining)
```

```
## [1] 19622 160
```

```
dim(pmlTesting)
```

```
## [1] 20 160
```

Data Cleansing

```
pmlTrainNoNA<-pmlTraining[, apply(pmlTraining, 2, function(x) !any(is.na(x)))]
dim(pmlTrainNoNA)
```

```
## [1] 19622 60
```

```
pmlTrainingClean<-pmlTrainNoNA[,-c(1:8)]
dim(pmlTrainingClean)
```

```
## [1] 19622 52
```

```
pmlTrainNzv <- nearZeroVar(pmlTrainingClean, saveMetrics=TRUE)
pmlTrainingCleanNew <- pmlTrainingClean[,pmlTrainNzv$nzv==FALSE]
dim(pmlTrainingCleanNew)
```

```
## [1] 19622    52
```

```
pmlTestingClean<-pmlTesting[,names(pmlTrainingCleanNew[,,-52])]
dim(pmlTestingClean)
```

```
## [1] 20 51
```

Data Partitioning and Prediction

```
inTrainingData<-createDataPartition(y=pmlTrainingCleanNew$classe, p=0.60,list=FALSE)
myTrainData <- pmlTrainingCleanNew[inTrainingData,]
dim(myTrainData)
```

```
## [1] 11776    52
```

```
myTestData <- pmlTrainingClean[-inTrainingData,]
dim(myTestData)
```

```
## [1] 7846    52
```

Results and Conclusions

```
set.seed(12345)
modFit1 <- rpart(classe ~ ., data=myTrainData, method="class")
fancyRpartPlot(modFit1)
```

```
## Warning: labs do not fit even at cex 0.15, there may be some overplotting
```


	Class: A	Class: B	Class: C	Class: D	Class: E
## Sensitivity	0.8754	0.49736	0.7251	0.7341	0.5811
## Specificity	0.9129	0.94817	0.9179	0.8852	0.9592
## Pos Pred Value	0.7998	0.69714	0.6509	0.5563	0.7625
## Neg Pred Value	0.9485	0.88718	0.9405	0.9444	0.9105
## Prevalence	0.2845	0.19347	0.1744	0.1639	0.1838
## Detection Rate	0.2490	0.09623	0.1264	0.1203	0.1068
## Detection Prevalence	0.3114	0.13803	0.1942	0.2163	0.1401
## Balanced Accuracy	0.8942	0.72277	0.8215	0.8096	0.7702

Prediction with Regression

```
fitControl1<-trainControl(method="cv", number=5, allowParallel=T, verbose=T)
gbmfit<-train(classe~.,data=myTrainData, method="gbm", trControl=fitControl1, verbose=F)
```

```
## Loading required package: gbm
```

```
## Warning: package 'gbm' was built under R version 3.2.5
```

```
## Loading required package: survival
```

```
##
```

```
## Attaching package: 'survival'
```

```
## The following object is masked from 'package:caret':
```

```
##
```

```
##      cluster
```

```
## Loading required package: splines
```

```
## Loading required package: parallel
```

```
## Loaded gbm 2.1.1
```

```
## Loading required package: plyr
```

```
## + Fold1: shrinkage=0.1, interaction.depth=1, n.minobsinnode=10, n.trees=150
## - Fold1: shrinkage=0.1, interaction.depth=1, n.minobsinnode=10, n.trees=150
## + Fold1: shrinkage=0.1, interaction.depth=2, n.minobsinnode=10, n.trees=150
## - Fold1: shrinkage=0.1, interaction.depth=2, n.minobsinnode=10, n.trees=150
## + Fold1: shrinkage=0.1, interaction.depth=3, n.minobsinnode=10, n.trees=150
## - Fold1: shrinkage=0.1, interaction.depth=3, n.minobsinnode=10, n.trees=150
## + Fold2: shrinkage=0.1, interaction.depth=1, n.minobsinnode=10, n.trees=150
## - Fold2: shrinkage=0.1, interaction.depth=1, n.minobsinnode=10, n.trees=150
## + Fold2: shrinkage=0.1, interaction.depth=2, n.minobsinnode=10, n.trees=150
## - Fold2: shrinkage=0.1, interaction.depth=2, n.minobsinnode=10, n.trees=150
## + Fold2: shrinkage=0.1, interaction.depth=3, n.minobsinnode=10, n.trees=150
## - Fold2: shrinkage=0.1, interaction.depth=3, n.minobsinnode=10, n.trees=150
## + Fold3: shrinkage=0.1, interaction.depth=1, n.minobsinnode=10, n.trees=150
## - Fold3: shrinkage=0.1, interaction.depth=1, n.minobsinnode=10, n.trees=150
## + Fold3: shrinkage=0.1, interaction.depth=2, n.minobsinnode=10, n.trees=150
```

```
## - Fold3: shrinkage=0.1, interaction.depth=2, n.minobsinnode=10, n.trees=150
## + Fold3: shrinkage=0.1, interaction.depth=3, n.minobsinnode=10, n.trees=150
## - Fold3: shrinkage=0.1, interaction.depth=3, n.minobsinnode=10, n.trees=150
## + Fold4: shrinkage=0.1, interaction.depth=1, n.minobsinnode=10, n.trees=150
## - Fold4: shrinkage=0.1, interaction.depth=1, n.minobsinnode=10, n.trees=150
## + Fold4: shrinkage=0.1, interaction.depth=2, n.minobsinnode=10, n.trees=150
## - Fold4: shrinkage=0.1, interaction.depth=2, n.minobsinnode=10, n.trees=150
## + Fold4: shrinkage=0.1, interaction.depth=3, n.minobsinnode=10, n.trees=150
## - Fold4: shrinkage=0.1, interaction.depth=3, n.minobsinnode=10, n.trees=150
## + Fold5: shrinkage=0.1, interaction.depth=1, n.minobsinnode=10, n.trees=150
## - Fold5: shrinkage=0.1, interaction.depth=1, n.minobsinnode=10, n.trees=150
## + Fold5: shrinkage=0.1, interaction.depth=2, n.minobsinnode=10, n.trees=150
## - Fold5: shrinkage=0.1, interaction.depth=2, n.minobsinnode=10, n.trees=150
## + Fold5: shrinkage=0.1, interaction.depth=3, n.minobsinnode=10, n.trees=150
## - Fold5: shrinkage=0.1, interaction.depth=3, n.minobsinnode=10, n.trees=150
## Aggregating results
## Selecting tuning parameters
## Fitting n.trees = 150, interaction.depth = 3, shrinkage = 0.1, n.minobsinnode = 10 on full training s
```

```
gbmfit$finalModel
```

```
## A gradient boosted model with multinomial loss function.
## 150 iterations were performed.
## There were 51 predictors of which 42 had non-zero influence.
```

```
class(gbmfit)
```

```
## [1] "train"          "train.formula"
```

```
predict2<-predict(gbmfit, newdata=myTrainData)
confusionMatrix(predict2, myTrainData$classe)
```

```
## Confusion Matrix and Statistics
```

```
##
##           Reference
## Prediction   A    B    C    D    E
##           A 3320   44    0    0    1
##           B   16 2189   46    0   16
##           C    11   44 1985   65   11
##           D     0    1   19 1856   16
##           E     1    1    4    9 2121
##
```

```
## Overall Statistics
```

```
##
##           Accuracy : 0.9741
##           95% CI : (0.9711, 0.9769)
##           No Information Rate : 0.2843
##           P-Value [Acc > NIR] : < 2.2e-16
##
##           Kappa : 0.9672
##           Mcnemar's Test P-Value : NA
##
```

```
## Statistics by Class:
##
##
##          Class: A Class: B Class: C Class: D Class: E
## Sensitivity      0.9916  0.9605  0.9664  0.9617  0.9797
## Specificity      0.9947  0.9918  0.9865  0.9963  0.9984
## Pos Pred Value   0.9866  0.9656  0.9381  0.9810  0.9930
## Neg Pred Value    0.9967  0.9905  0.9929  0.9925  0.9954
## Prevalence       0.2843  0.1935  0.1744  0.1639  0.1838
## Detection Rate    0.2819  0.1859  0.1686  0.1576  0.1801
## Detection Prevalence 0.2858  0.1925  0.1797  0.1607  0.1814
## Balanced Accuracy 0.9931  0.9761  0.9765  0.9790  0.9891
```

```
predgbm<-predict(gbmfit, newdata=myTestData)
confusionMatrix(predgbm, myTestData$classe)
```

```
## Confusion Matrix and Statistics
##
##          Reference
## Prediction    A    B    C    D    E
##          A 2189   48    0    1    2
##          B   31 1419   49    4   16
##          C   10   41 1306   49   15
##          D    1    6   12 1223   17
##          E    1    4    1    9 1392
##
## Overall Statistics
##
##          Accuracy : 0.9596
##          95% CI : (0.955, 0.9638)
##          No Information Rate : 0.2845
##          P-Value [Acc > NIR] : < 2.2e-16
##
##          Kappa : 0.9489
##          Mcnemar's Test P-Value : 4.591e-09
##
```

```
## Statistics by Class:
##
##          Class: A Class: B Class: C Class: D Class: E
## Sensitivity      0.9807  0.9348  0.9547  0.9510  0.9653
## Specificity      0.9909  0.9842  0.9822  0.9945  0.9977
## Pos Pred Value   0.9772  0.9342  0.9191  0.9714  0.9893
## Neg Pred Value    0.9923  0.9844  0.9904  0.9904  0.9922
## Prevalence       0.2845  0.1935  0.1744  0.1639  0.1838
## Detection Rate    0.2790  0.1809  0.1665  0.1559  0.1774
## Detection Prevalence 0.2855  0.1936  0.1811  0.1605  0.1793
## Balanced Accuracy 0.9858  0.9595  0.9685  0.9728  0.9815
```

Prediction with Random Forests

```
fitControl <- trainControl(method="cv", number=5, allowParallel=T, verbose=T)
rffit<-train(classe~.,data=myTrainData, method="rf", trControl=fitControl, verbose=F)
```

```
## + Fold1: mtry= 2
```

```

## - Fold1: mtry= 2
## + Fold1: mtry=26
## - Fold1: mtry=26
## + Fold1: mtry=51
## - Fold1: mtry=51
## + Fold2: mtry= 2
## - Fold2: mtry= 2
## + Fold2: mtry=26
## - Fold2: mtry=26
## + Fold2: mtry=51
## - Fold2: mtry=51
## + Fold3: mtry= 2
## - Fold3: mtry= 2
## + Fold3: mtry=26
## - Fold3: mtry=26
## + Fold3: mtry=51
## - Fold3: mtry=51
## + Fold4: mtry= 2
## - Fold4: mtry= 2
## + Fold4: mtry=26
## - Fold4: mtry=26
## + Fold4: mtry=51
## - Fold4: mtry=51
## + Fold5: mtry= 2
## - Fold5: mtry= 2
## + Fold5: mtry=26
## - Fold5: mtry=26
## + Fold5: mtry=51
## - Fold5: mtry=51
## Aggregating results
## Selecting tuning parameters
## Fitting mtry = 26 on full training set

```

```

predict3<-predict(rffit, newdata=myTrainData)
confusionMatrix(predict3, myTrainData$classe)

```

```

## Confusion Matrix and Statistics
##
##           Reference
## Prediction    A    B    C    D    E
##           A 3348    0    0    0    0
##           B    0 2279    0    0    0
##           C    0    0 2054    0    0
##           D    0    0    0 1930    0
##           E    0    0    0    0 2165
##
## Overall Statistics
##
##           Accuracy : 1
##           95% CI : (0.9997, 1)
##           No Information Rate : 0.2843
##           P-Value [Acc > NIR] : < 2.2e-16
##
##           Kappa : 1

```



```
## McNemar's Test P-Value : NA
##
## Statistics by Class:
##
##           Class: A Class: B Class: C Class: D Class: E
## Sensitivity      1.0000  1.0000  1.0000  1.0000  1.0000
## Specificity      1.0000  1.0000  1.0000  1.0000  1.0000
## Pos Pred Value   1.0000  1.0000  1.0000  1.0000  1.0000
## Neg Pred Value   1.0000  1.0000  1.0000  1.0000  1.0000
## Prevalence       0.2843  0.1935  0.1744  0.1639  0.1838
## Detection Rate   0.2843  0.1935  0.1744  0.1639  0.1838
## Detection Prevalence 0.2843  0.1935  0.1744  0.1639  0.1838
## Balanced Accuracy 1.0000  1.0000  1.0000  1.0000  1.0000
```

```
predrf<-predict(rffit, newdata=myTestData)
confusionMatrix(predrf, myTestData$classe)
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction    A    B    C    D    E
##           A 2226   14    0    0    0
##           B    5 1492   11    2    0
##           C    0   10 1352    8    3
##           D    0    0    5 1275    6
##           E    1    2    0    1 1433
##
## Overall Statistics
##
##           Accuracy : 0.9913
##           95% CI : (0.989, 0.9933)
##           No Information Rate : 0.2845
##           P-Value [Acc > NIR] : < 2.2e-16
##
##           Kappa : 0.989
## McNemar's Test P-Value : NA
##
## Statistics by Class:
##
##           Class: A Class: B Class: C Class: D Class: E
## Sensitivity      0.9973  0.9829  0.9883  0.9914  0.9938
## Specificity      0.9975  0.9972  0.9968  0.9983  0.9994
## Pos Pred Value   0.9938  0.9881  0.9847  0.9914  0.9972
## Neg Pred Value   0.9989  0.9959  0.9975  0.9983  0.9986
## Prevalence       0.2845  0.1935  0.1744  0.1639  0.1838
## Detection Rate   0.2837  0.1902  0.1723  0.1625  0.1826
## Detection Prevalence 0.2855  0.1925  0.1750  0.1639  0.1832
## Balanced Accuracy 0.9974  0.9900  0.9925  0.9949  0.9966
```

```
predictpmlTesting<-predict(rffit, newdata=pmlTesting)
```

Output for the prediction of the 20 cases provided

```
predictpmlTesting
```

```
## [1] B A B A A E D B A A B C B A E E A B B B  
## Levels: A B C D E
```

```
pml_write_files = function(x){  
  n = length(x)  
  for(i in 1:20){  
    filename = paste0("problem_id_",i,".txt")  
    write.table(x[i],file=filename,quote=FALSE,row.names=FALSE,col.names=FALSE)  
  }  
}  
pml_write_files(predictpmlTesting)
```